

Chapter 6

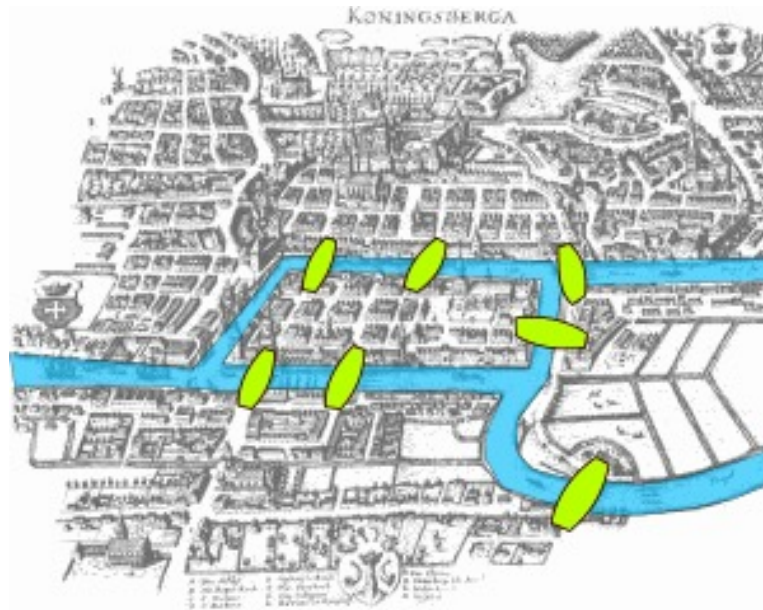
Introduction to Graphs

Topics to be Covered


- Learn about graphs
- Become familiar with the basic terminology of graph theory
- Discover how to represent graphs in computer memory
- Examine and implement various graph traversal algorithms
- Examine and implement some applications of graph theory

Introduction

- In 1736, in the town of Königsberg in Prussia, the river Pregel flows around the island Kneiphof and then divides into two branches



Introduction

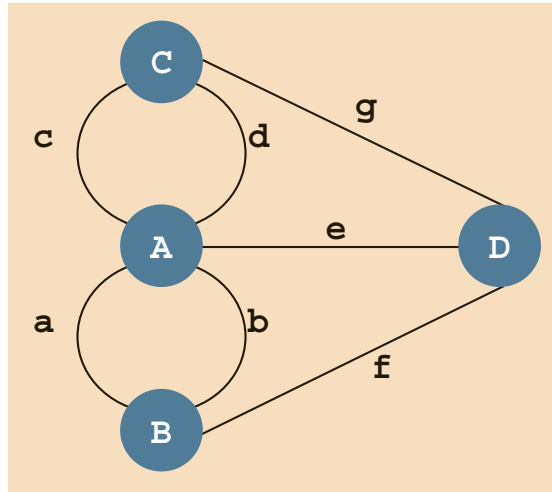


- The river has four land areas: A, B, C, and D
- These land areas are connected using seven bridges that are labeled a , b , c , d , e , f , and g
- *The Königsberg bridge problem*

Starting at one land area, is it possible to walk across all of the bridges exactly once and return to the starting land area?

Introduction

- In 1736, Euler represented the *Königsberg bridge problem* as a *graph* ...



... and answered the question in the negative

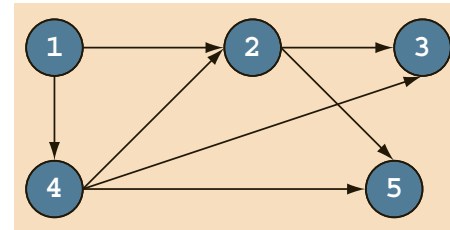
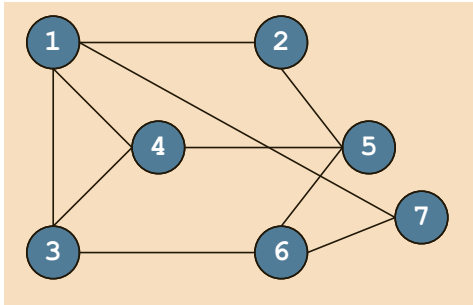
➡ This marked the birth of *graph theory*

Graph Definitions and Notations

- A *graph* G is a pair, $G = (V, E)$, in which V is the set of vertices of G , and E is the set of edges in G
- If the elements of E are ordered pairs, G is called a *directed graph* (or *digraph*); otherwise, G is called an *undirected graph*
 - In an undirected graph, the pairs (u, v) and (v, u) represent the same edge
 - If (u, v) is an edge in a digraph, the vertex u is called the *origin* of the edge, and the vertex v is called the *destination*

Graph Definitions and Notations

- A graph can be shown pictorially



- A graph is called a *weighted graph* if its edges are labeled with numeric values
- A graph $H = (V_H, E_H)$ is called a *subgraph* of G if

$$V_H \subseteq V, E_H \subseteq E$$

Graph Definitions and Notations

Let G be an undirected graph; u, v be two vertices of G

- u and v are called *adjacent* if there is an edge from one to the other; that is, $(u, v) \in E$
- Let $e = (u, v) \in E$. Edge e is *incident on* the vertices u and v
- An edge incident on a single vertex is called a *loop*
- If two edges, e_1 and e_2 , are associated with the same pair of vertices, then e_1 and e_2 are called *parallel edges*

Graph Definitions and Notations

- A graph is called a *simple graph* if it has no loops and no parallel edges
- There is a *path* from u to v if there is a sequence of vertices u_1, u_2, \dots, u_n such that $u = u_1, v = u_n$, and (u_i, u_{i+1}) is an edge for all $i = 1, 2, \dots, n - 1$
- A *simple path* from u to v is a path from u to v with no repeated vertices
- A *cycle* is a simple path in which the first and last vertices are the same

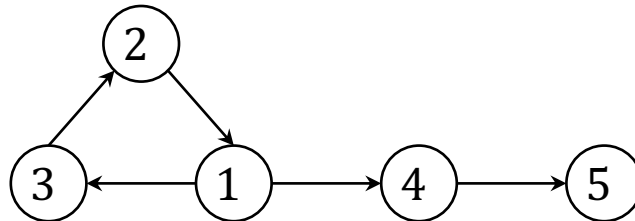
Graph Definitions and Notations

- Vertices u and v are called *connected* if there is a path from u to v
- G is called *connected* if there is a path from any vertex to any other vertex
 - A subset of connected vertices is called a *connected component* of G

Graph Definitions and Notations

Let G be a directed graph; u, v be two vertices in G

- If $(u, v) \in E$ then we say that u is *adjacent to* v and v is *adjacent from* u
- G is called *strongly connected* if any two vertices in G are connected
 - A *strongly connected component* of G is a maximal strongly connected subgraph



Graph Definitions and Notations

Let G be a directed graph; u, v be two vertices in G

- If $(u, v) \in E$ then we say that u is *adjacent to* v and v is *adjacent from* u
- G is called *strongly connected* if any two vertices in G are connected
 - A *strongly connected component* of G is a maximal strongly connected subgraph
- The *outdegree* of v is the number of directed edges leaving v ; the *indegree* of v is the number of directed edges entering v

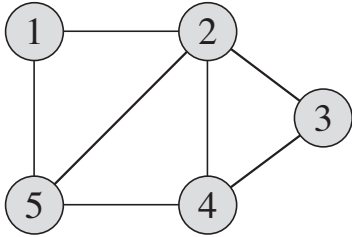
Graph Representation

- A graph can be represented in several ways
 - How a graph is represented in memory depends on the specific application
- The most two common methods are *adjacency matrices* and *adjacency lists*
- Let $G = (V, E)$ be a graph with $n(= |V|)$ vertices
 - Let $V = \{v_1, v_2, \dots, v_n\}$

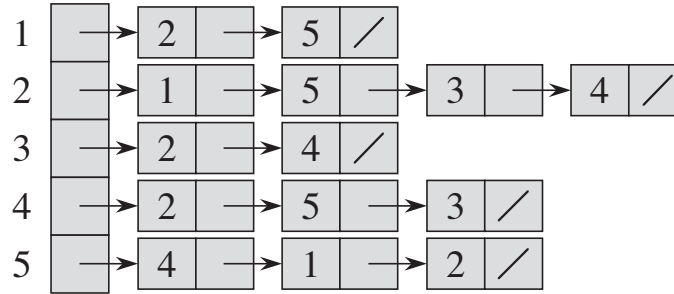
Graph Representation: Adjacency Matrices

- The *adjacency matrix* A of G is a two-dimensional $n \times n$ matrix such that
 - If there is an edge from v_i to v_j , the $(i, j)^{th}$ entry of A is 1
 - Otherwise, the $(i, j)^{th}$ entry is 0
- In an undirected graph, if $(v_i, v_j) \in E$ then $(v_j, v_i) \in E$, so the $(i, j)^{th}$ entry is as same as the $(j, i)^{th}$ entry
 - ➡ The adjacency matrix of an undirected graph is symmetric

Graph Representation: Examples



(a)

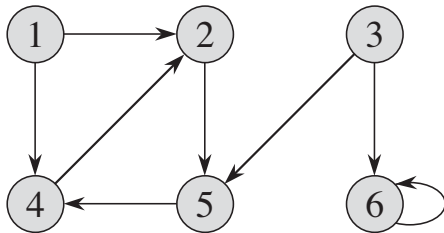


(b)

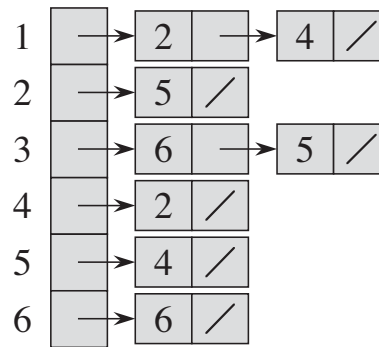
	1	2	3	4	5
1	0	1	0	0	1
2	1	0	1	1	1
3	0	1	0	1	0
4	0	1	1	0	1
5	1	1	0	1	0

(c)

Two representations of an undirected graph



(a)



(b)

	1	2	3	4	5	6
1	0	1	0	1	0	0
2	0	0	0	0	1	0
3	0	0	0	0	1	1
4	0	1	0	0	0	0
5	0	0	0	1	0	0
6	0	0	0	0	0	1

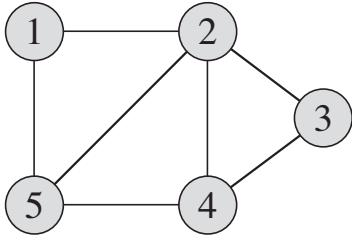
(c)

Two representations of a directed graph

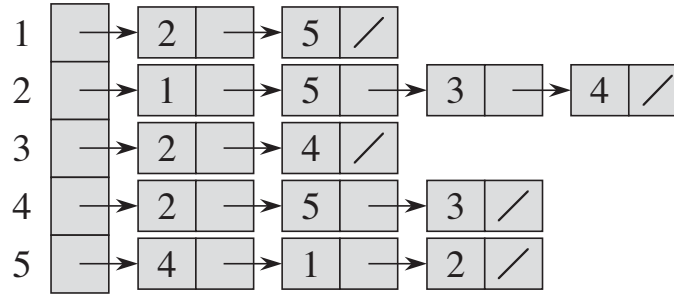
Graph Representation: Adjacency Lists

- With *adjacency lists*, corresponding to each vertex, u , there is a linked list such that each node of the linked list contains the vertex, v , such that $(u, v) \in E$
- Technically, we use an array A of size n , such that $A[i]$ is
 - a representation of the vertex v_i
 - a *pointer* to the first node of the linked list containing the vertices *to which v_i is adjacent*

Graph Representation: Examples



(a)

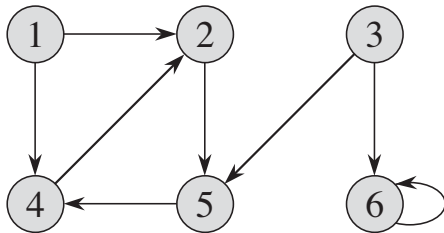


(b)

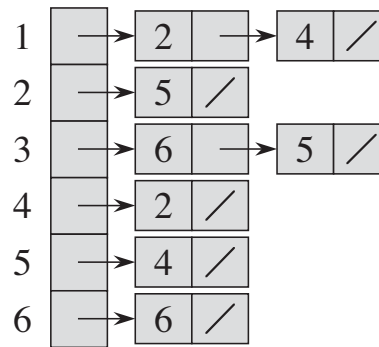
	1	2	3	4	5
1	0	1	0	0	1
2	1	0	1	1	1
3	0	1	0	1	0
4	0	1	1	0	1
5	1	1	0	1	0

(c)

Two representations of an undirected graph



(a)



(b)

	1	2	3	4	5	6
1	0	1	0	1	0	0
2	0	0	0	0	1	0
3	0	0	0	0	1	1
4	0	1	0	0	0	0
5	0	0	0	1	0	0
6	0	0	0	0	0	1

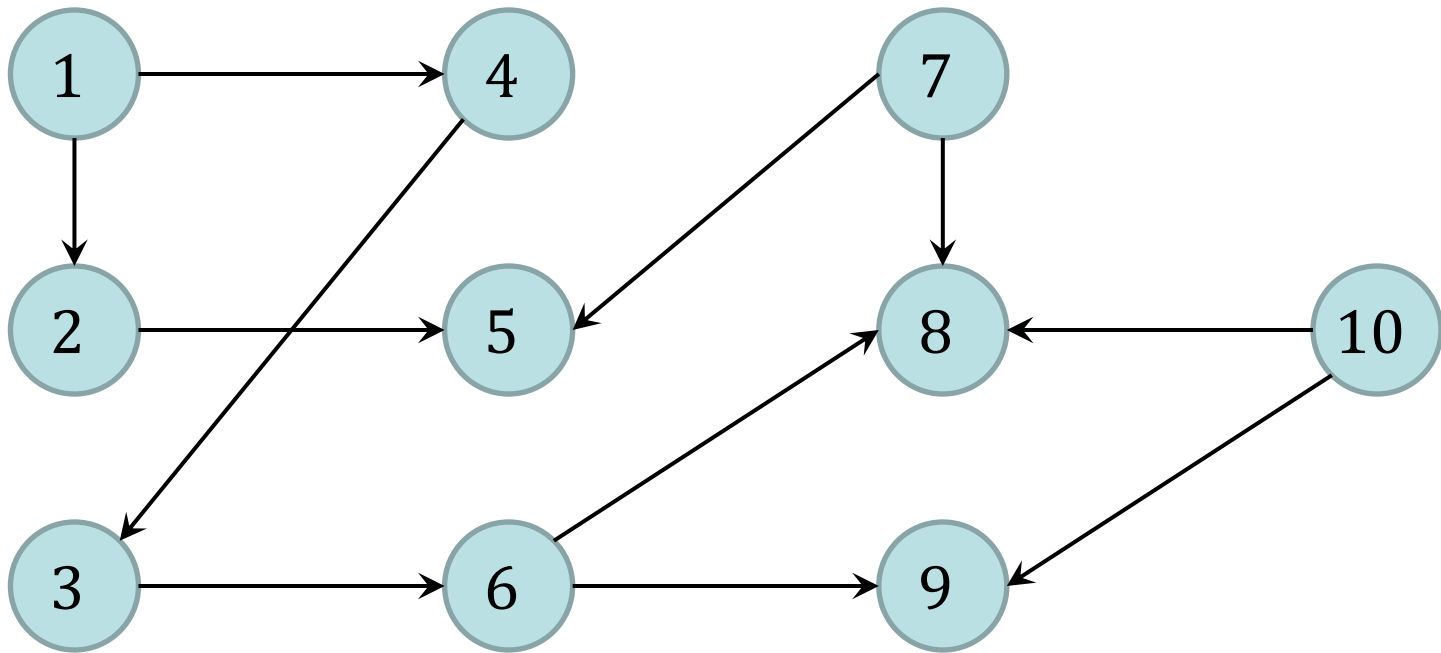
(c)

Two representations of a directed graph

Graph Traversals

- Traversing a graph is a bit more complicated than traversing a binary tree
 - A binary tree has no cycles, but a graph might have cycles
 - We might not be able to traverse the entire graph from a single vertex
- In order to traverse the entire graph, we must ...
 - keep track of the vertices that have been visited
 - traverse the graph from each unvisited vertex

Depth First Traversal



A depth first ordering of the vertices of the graph is

1, 2, 5, 4, 3, 6, 8, 9, 7, 10

Depth First Traversal: A Non-recursive Algorithm

Mark each vertex in V as unvisited

for each vertex v in V

if v is an unvisited vertex

Push v onto the empty stack

while the stack is not empty

Pop u off the stack

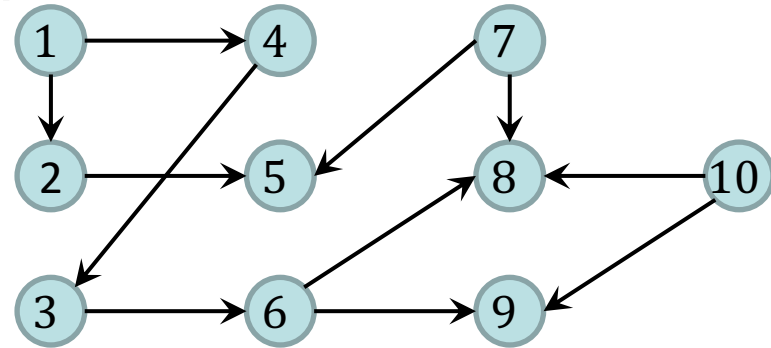
if u is an unvisited vertex

Visit u and mark it as visited

for each vertex w to which u is adjacent

if w is an unvisited vertex

Push w onto the stack



Depth First Traversal: Pseudo-code

```
dft(graph) {  
    visited[1 .. n] = false;  
    for (v = 1; v ≤ n; v++)  
        if (!visited[v]) {  
            push(S, v);  
            while (!isEmpty(S)) {  
                u = pop(S);  
                if (!visited[u]) {  
                    cout << u;  
                    visited[u] = true;  
                    for (p = graph[u]; p; p = p->next)  
                        if (!visited[p->vertex])  
                            push(S, p->vertex);  
                }  
            }  
        }  
}
```

Depth First Traversal: A Recursive Algorithm

$\text{dft}(v)$ // *DFT starts at the vertex v*

Visit v and mark it as visited

for each vertex u to which v is adjacent

if u is an unvisited vertex

$\text{dft}(u)$

Mark each vertex in V as unvisited

for each vertex v in V

if v is an unvisited vertex

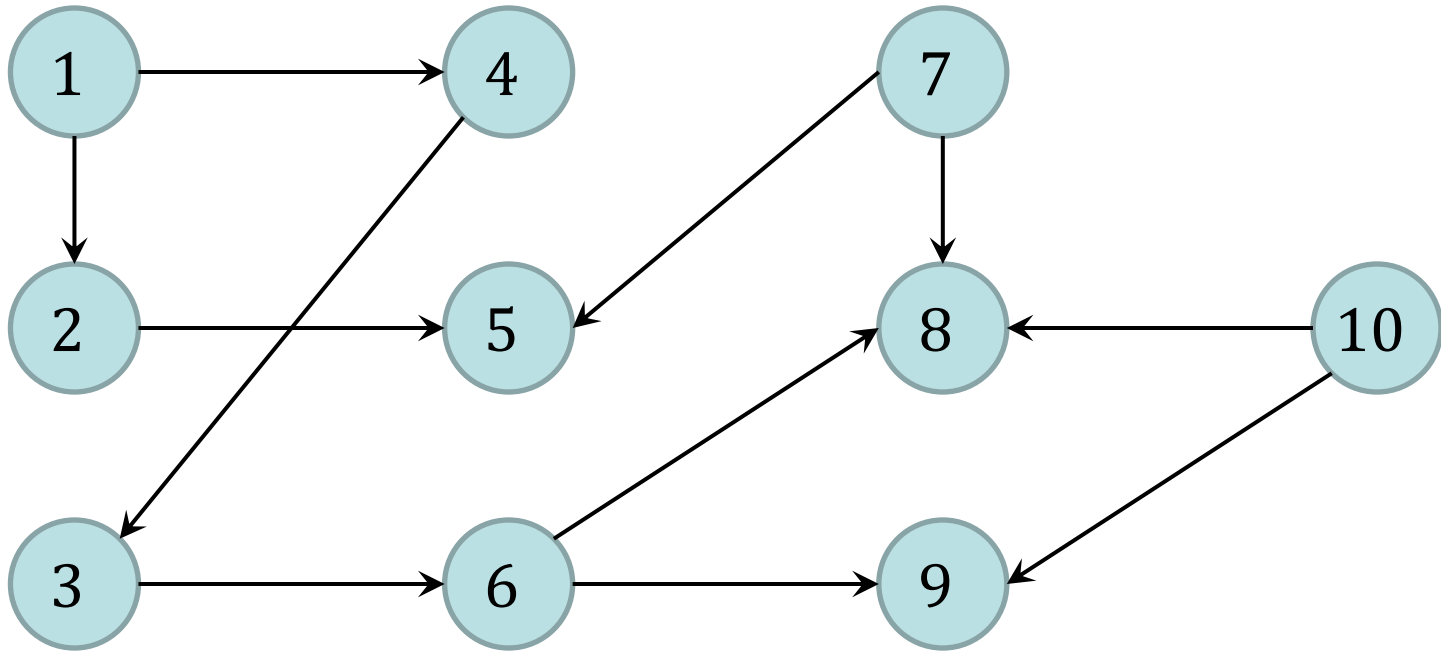
$\text{dft}(v)$

Depth First Traversal: Pseudo-code

```
dft(graph, v, visited) {  
    visited[v] = true;  
    cout << v;  
    for (p = graph[v]; p; p = p->next;)  
        if (!visited[p->vertex])  
            dft(graph, p->vertex, visited);  
}
```

```
visited[1 .. n] = false;  
for (v = 1; v ≤ n; v++)  
    if (!visited[v])  
        dft(graph, v, visited);
```

Breadth First Traversal



A breadth first ordering of the vertices of the graph is

1, 2, 4, 5, 3, 6, 8, 9, 7, 10

Breadth First Traversal: Algorithm

Mark each vertex in V as unvisited

for each vertex v in the graph

if v is an unvisited vertex

 Add v to the queue

 Mark v as visited

while the queue is not empty

 Extract u from the queue

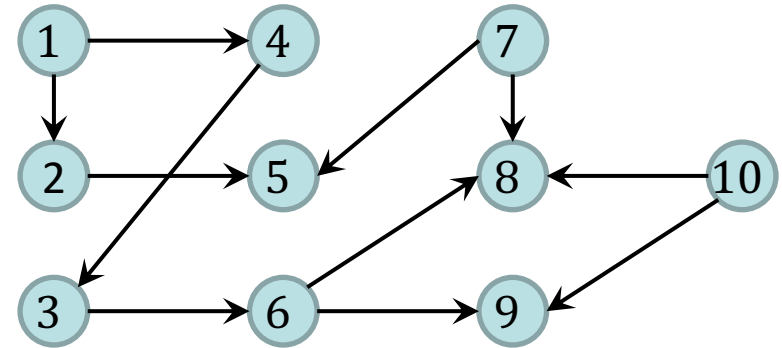
 Visit u

for each vertex w to which u is adjacent

if w is an unvisited vertex

 Add w to the queue

 Mark w as visited



Breadth First Traversal: Pseudo-code

```
bft(graph) {  
    visited[1 .. n] = false;  
    for (v = 1; v ≤ n; v++)  
        if (!visited[v]) {  
            enqueue(Q, v);  
            visited[v] = true;  
            while (!isEmpty(Q)) {  
                u = dequeue(Q);    cout << u;  
                for (p = graph[u]; p; p = p->next)  
                    if (!visited[p->vertex]) {  
                        enqueue(Q, p->vertex);  
                        visited[p->vertex] = true;  
                    }  
            }  
        }  
}
```

Topological Sorting

Linear ordering vs. Partial ordering

- A *linear ordering* on a finite set of items is an ordering which is given over *all pairs of items*
- A *partial ordering* on a finite set of items is an ordering which is given over *some pairs of items* but not among all of them



Partial Ordering: Example

- A task is broken up into subtasks and completion of certain subtasks must usually precede the execution of other subtasks
 - If a subtask v must precede a subtask w : $v \angle w$
- In a curriculum of computer science, certain courses must be taken before others
 - If course v is a prerequisite for course w : $v \angle w$

Topological Sorting: Properties



- A partial ordering on a set S satisfies the following properties for any distinct items x , y , and z of S :
 - Transitivity: If $x \angle y$ and $y \angle z$, then $x \angle z$
 - Asymmetry: If $x \angle y$, then not $y \angle x$
 - Irreflexivity: Not $x \angle x$
- Hence, a partial ordering can be illustrated by drawing a *directed acyclic graph* (DAG) in which ...
 - the vertices denote the items of S
 - the arrows represent ordering relationships

DAGs



Topological Sorting

- The problem of topological sorting is to embed the partial order in a linear order
 - Graphically, if the vertices are arranged linearly and in a *topological order*, the arrows will all point in one direction
- The vertices in a DAG may have several topological orders



Topological Sorting: Algorithm

Step 1: Add all vertices whose *indegree* is 0 to a queue

Step 2: Do the following substeps repeatedly until the queue is empty

Step 2.1: Take a vertex v off the queue and add v to the *end* of the resulting list

Step 2.2: Remove vertex v and the edges that *leave* it from the graph

Step 2.3: Some vertices whose *indegree* is 0 may occur in the graph after *Step 2.2*. Add them to the queue

Topological Sorting: An Alternative Version

Step 1: Add all vertices whose *outdegree* is 0 to a queue

Step 2: Do the following substeps repeatedly until the queue is empty

Step 2.1: Take a vertex v off the queue and add v to the *beginning* of the resulting list


Step 2.2: Remove vertex v and the edges that *enter* it from the graph

Step 2.3: Some vertices whose *outdegree* is 0 may occur in the graph after *Step 2.2*. Add them to the queue

Spanning Trees


- A tree is a special kind of undirected graph, one that is connected but that has no cycles
 - Although all trees are graphs, not all graphs are trees
- A *spanning tree* of a connected undirected graph G is a subgraph of G that contains all of G 's vertices and enough of its edges to form a tree
- There may be several spanning trees for a given graph

Spanning Trees: Example



Spanning Trees: Some Observations

- A connected undirected graph that has n vertices ...
 - ... must have at least $n - 1$ edges
 - ... and exactly $n - 1$ edges cannot contain a cycle
 - ... and more than $n - 1$ edges must contain at least one cycle

 To obtain the spanning tree of a connected graph of n vertices, we must connect its n vertices with $n - 1$ edges

Spanning Trees: Algorithm

- Beginning at a vertex, our computer visits all other vertices in the graph
 - Each vertex will only be visited once
- As our computer traverses the graph, it also marks the edge that it follows
- After the traversal is complete, the graph's vertices and marked edges form a spanning tree
 - The unmarked edges can be removed from the graph

The DFS Spanning Tree

dfsTree(v)

Mark v as visited

for each vertex u to which v is adjacent

if u is an unvisited vertex

Mark the edge between v and u

dfsTree(u)

The BFS Spanning Tree

bfsTree(v)

Add v to an empty queue

Mark v as visited

while the queue is not empty

 Extract u from the queue

for each vertex w to which u is adjacent

if w is an unvisited vertex

 Mark edge between u and w


 Add w to the queue

 Mark w as visited


Minimum Spanning Trees

- A *minimum spanning tree* (or *mst*) of an undirected weighted connected graph is its spanning tree of the smallest weight
 - The *weight of a tree* is defined as the sum of the weights on all its edges
- It has direct applications to the design of all kinds of networks by providing the cheapest way to achieve connectivity


Example




Graph



$$w(T_1) = 6$$



$$w(T_2) = 9$$



$$w(T_3) = 8$$

Minimum Spanning Trees: Prim's Algorithm

- The algorithm constructs a *mst* through a sequence of expanding subtrees:
 - The initial subtree consists of a single vertex selected arbitrarily from the set V
 - On each iteration, the algorithm expands the current tree by attaching to it the *nearest vertex* not in that tree
 - The algorithm stops after all the graph's vertices have been included in the tree being constructed
- ➡ The total number of such iterations is $|V| - 1$

An Outline of Prim's Algorithm

Prim($G = (V, T)$)

$$V_T = \{v_1\}$$

$$E_T = \emptyset$$

for $i = 1$ to $|V| - 1$

Find a minimum-weight edge $e^* = (u^*, v^*)$ among all edges (u, v) such that $u \in V_T, v \in V \setminus V_T$

$$V_T = V_T \cup \{v^*\}$$

$$E_T = E_T \cup \{e^*\}$$

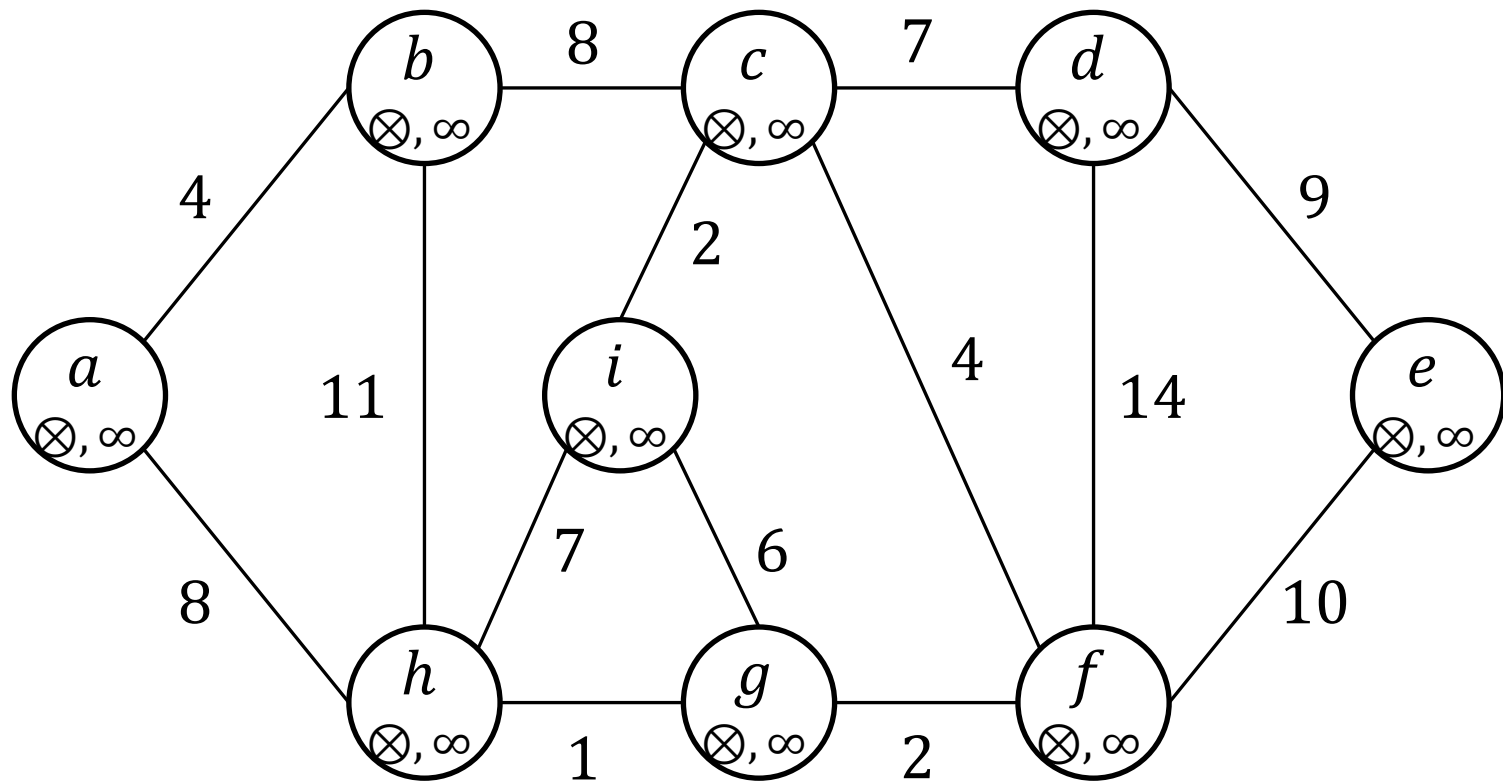
return $G_T = (V_T, E_T)$

An Implementation of Prim's Algorithm

- Each vertex in $V \setminus V_T$ needs to remember the shortest edge connecting the vertex to a vertex in V_T
- Each vertex has two labels:
 - “parent” label: The name of the nearest tree vertex ($\in V_T$)
 - “distance” label: The weight of the corresponding edge
- A vertex that is not adjacent to any of the tree vertices:
 - “parent” label is NIL
 - “distance” label is ∞

An Implementation of Prim's Algorithm

- A vertex ($\in V \setminus V_T$) with the smallest distance label is the next one to be added to the tree being constructed G_T
 - Ties can be broken arbitrarily
- Let v^* be the next vertex to be added to G_T
 - Move v^* from the set $V \setminus V_T$ to the set V_T
 - For each remaining vertex v in $V \setminus V_T$ that is connected to v^* by a shorter edge than the v 's current distance label, update its labels by v^* and $G[v^*][v]$, respectively



```

Prim(G, root) {
    for (each vertex  $v \in V$ ) {
         $v.dist = \infty$ ;
         $v.parent = NIL$ ;
    }
     $root.dist = 0$ ;       $V_T = \emptyset$ ;
    createQueue(pQ, V);
    while (!isEmpty(pQ)) {
         $v^* = extractQueue(pQ)$ ;
        add  $v^*$  to  $V_T$ ;
        for (each  $v \in pQ$  that is adjacent to  $v^*$ )
            if ( $G[v^*][v] < v.dist$ ) {
                 $v.dist = G[v^*][v]$ ;
                 $v.parent = v^*$ ;
                updateQueue(pQ, v);
            }
    }
}

```

Single-Source Shortest-Paths Problem

- For a given vertex called the *source* in a weighted connected graph, find shortest paths to all its other vertices
- The most widely applications of the problem are ...
 - transportation planning
 - packet routing in communication networks, including the Internet

Dijkstra's Algorithm: The General Idea

- The algorithm finds the shortest path from the *source* to a vertex nearest to it, then to a second nearest, ...
- ➡ Before the i^{th} iteration starts, the algorithm has already identified the shortest paths to $i - 1$ other vertices nearest to the *source*
 - These vertices, the *source*, and the edges of the shortest paths leading to them from the *source* form a subtree T_i
- The next vertex nearest to the *source* can be found among the vertices adjacent to the vertices of T_i

An Outline of Dijkstra's Algorithm

// d_u the length of the shortest path from the *source* to u

Dijkstra($G = (V, T)$)

$V_T = \{v_1\}$ // v_1 is the *source*

$E_T = \emptyset$

for $i = 1$ to $|V| - 1$

Find an edge $e^* = (u^*, v^*)$ among all edges (u, v)

such that $u \in V_T, v \in V \setminus V_T$ and $d_{u^*} + G[u^*][v^*] \leq d_u + G[u][v]$

$V_T = V_T \cup \{v^*\}$

$E_T = E_T \cup \{e^*\}$


return $G_T = (V_T, E_T)$

An Implementation of Dijkstra's Algorithm

- Each vertex has two labels:
 - “parent” label: The name of the nearest tree vertex ($\in V_T$)
 - “distance” label: The length of the shortest path from the source to this vertex found by the algorithm so far
- ➡ When a vertex is added to V_T , this label indicates the length of the shortest path from the *source* to that vertex
- A vertex that is not adjacent to any of the tree vertices:
 - “parent” label is NIL
 - “distance” label is ∞

An Implementation of Dijkstra's Algorithm

- Let v^* be the next vertex to be added to G_T
 - Move v^* from the set $V \setminus V_T$ to the set V_T
 - For each remaining vertex v in $V \setminus V_T$ that is connected to v^* such that $d_{v^*} + G[v^*][v] < d_v$, update the labels of v by v^* and $d_{v^*} + G[v^*][v]$, respectively



```

Dijkstra( $G(V, E)$ , source) {
  for (each vertex  $v \in V$ ) {
     $v.dist = \infty$ ;
     $v.parent = NIL$ ;
  }
   $source.dist = 0$ ;  $V_T = \emptyset$ ;
  createQueue(pQ, V);
  while (!isEmpty(pQ)) {
     $v^* = extractQueue(pQ)$ ;
    add  $v^*$  to  $V_T$ ;
    for (each  $v \in pQ$  that is adjacent to  $v^*$ )
      if ( $v^*.dist + G[v^*][v] < v.dist$ ) {
         $v.parent = v^*$ ;
         $v.dist = v^*.dist + G[v^*][v]$ ;
        updateQueue(pQ, v);
      }
  }
}

```