

# DATA STRUCTURES & ALGORITHMS

## Lecture 6: TREES – Part 2


### Balanced Binary Search Tree

Lecturer: Dr. Nguyen Hai Minh

# CONTENT


---

- Introduction
- AVL Trees
- Red-Black Trees



# Introduction

- **Balanced search tree**: A search-tree data structure for which a height of  $O(\log_2 n)$  is guaranteed when implementing a dynamic set of  $n$  items.
- Examples:
  - AVL trees
  - 2-3 trees
  - 2-3-4 trees
  - B-trees
  - Red-black trees




# AVL TREES

- The Tree ADT
- Tree Traversal




# AVL trees

- Proposed by **A**del'son-**V**el'skii and **L**andis in 1962.
- An AVL tree (originally called an admissible tree) is one in which the height of the left and right subtrees of every node differ by at most one.



# AVL trees

□ Example:



□ Each node has a *balance factor* that is the difference between the heights of the left and right subtrees.

- A balance factor is the height of the right subtree minus the height of the left subtree.
- Values: 0, -1, +1

# AVL trees

- Minimum number of nodes in an AVL tree:

$$AVL_h = AVL_{h-1} + AVL_{h-2} + 1$$

- $AVL_0 = 1$

- $AVL_1 = 2$

- Height of an AVL tree:  $O(\log_2 n)$


$$\log_2(n + 1) \leq h < 1.44\log_2(n + 2) - 0.328$$

- **Corollary.** The queries SEARCH, MIN, MAX, SUCCESSOR, and PREDECESSOR all run in  $O(\log_2 n)$  time on an AVL tree with  $n$  nodes.



# Balancing an AVL tree


- If balance factor of a node  $x$   $<-1$  or  $>+1$  (when INSERT or DELETE a node from an AVL tree): the tree needs to be balanced.
  - By “**rotation**” (*re-balance the subtree rooted with  $x$* )






# Balancing an AVL tree

- Let  $x$  be the unbalanced node
- 4 cases:
  1. Left-left unbalanced (LL): RIGHT-ROTATE( $x$ )
  2. Left-right unbalanced (LR): LEFT-ROTATE( $x.left$ ), then RIGHT-ROTATE( $x$ )
  3. Right-right unbalanced (RR): LEFT-ROTATE( $x$ )
  4. Right-left unbalanced (RL): RIGHT-ROTATE( $x.right$ ), then LEFT-ROTATE( $x$ )



# Rotations

- Rotations maintain the inorder ordering of keys:
  - $a \in \alpha, b \in \beta, c \in \gamma \Rightarrow a \leq A \leq b \leq B \leq c$ .
- A rotation can be performed in  $O(1)$  time.




# Insertion into an AVL tree

- The INSERT operation is performed in the same way with insertion in a binary search tree.
- If INSERT results in an unbalanced tree, perform the appropriate rotation(s) to restore its balance.




# Case 1: Left-left Unbalanced

□ Result of insert 1 into the tree



# Case 1: Left-left Unbalanced


## □ RIGHT-ROTATE(6)



## □ Even better than the original tree!


# Case 2: Left-right Unbalanced

□ Result of insert 4 into the tree




# Case 2: Left-right Unbalanced

□ LEFT-ROTATE(3) → Transformed to case 1 (LL)



# Case 2: Left-right Unbalanced


## □ RIGHT-ROTATE(6)






# Case 3: Right-right Unbalanced

□ Result of insert **25** into the tree



# Case 3: Right-right Unbalanced


## □ LEFT-ROTATE(17)



□ Even better than the original tree!


# Case 4: Right-left Unbalanced

□ Result of insert 20 into the tree




# Case 4: Right-left Unbalanced

□ RIGHT-ROTATE(23) → Transformed to case 3 (RR)



# Case 4: Right-left Unbalanced

□ LEFT-ROTATE(17)




# Deletion from an AVL tree

- The DELETE operation is performed in the same way with deletion in a binary search tree.
- After a node  $x$  is deleted:
  - Balance factors are updated from the parent of  $x$  up to the root.
  - For **each node in this path** whose balance factor becomes  $+2/-2$ , perform the appropriate rotation(s) to restore the balance of the tree.
  - Notice that the rebalancing **does not stop** after the first unbalanced node is rotated.



# Deletion without balancing


□ Delete 5 → Replace it by 8



□ The tree is still in balance → no rotation

# Deletion that needs balancing


□ Delete 3 → RL unbalanced → case 4 (RL)






# Deletion that needs balancing

□ RIGHT-ROTATE(10), then LEFT-ROTATE(5)



# Worst case of deletion

□ What happens when we delete 3?



# Analysis

## □ INSERT:

- Find a place to insert:  $O(h)$ .
- Perform 1 or 2 rotations.

$$O(\log_2 n)$$

## □ DELETE:

- Find a replaced node:  $O(h)$ .
- Perform rotations on the path from the deleted node up to the root:  $O(h)$ .

$$O(\log_2 n)$$




# RED-BLACK TREES

# Red-black trees


- Red-black tree (proposed by Rudolf Bayer in 1972) is a kind of self-balancing binary search tree. Each node of the tree has an extra bit as color field.
- Red-black properties:
  1. Every node is either **red** or black.
  2. The root is black.
  3. The leaves (**NIL**'s) are black.
  4. If a node is **red**, then both of its children are black.
  5. All simple paths from any node **x** to a descendent leaf have the same number of black nodes = **black-height(x)**



# Example of a red-black tree




# Example of a red-black tree



1. Every node is either red or black.


# Example of a red-black tree



2.+3. The root and leaves (NIL's) are black.




# Example of a red-black tree



4. If a node is **red**, then both its children are black.

# Example of a red-black tree



5. All simple paths from any node  $x$  to a descendent leaf have the same number of black nodes =  $black-height(x)$


# Height of a red-black tree

□ **Theorem.** A red-black tree with  $n$  keys has height

$$h \leq 2 \log_2(n+1)$$

□ **Proof. INTUITION:**

- Merge red nodes into their black parents




# Height of a red-black tree

□ **Theorem.** A red-black tree with  $n$  keys has height

$$h \leq 2 \log_2(n+1)$$

□ **Proof. INTUITION:**

- Merge their black parents




# Height of a red-black tree

□ **Theorem.** A red-black tree with  $n$  keys has height

$$h \leq 2 \log_2(n+1)$$

□ **Proof. INTUITION:**

- Merge their black parents




# Height of a red-black tree

□ **Theorem.** A red-black tree with  $n$  keys has height

$$h \leq 2 \log_2(n+1)$$

□ **Proof. INTUITION:**

- Merge red nodes into their black parents




# Height of a red-black tree

□ **Theorem.** A red-black tree with  $n$  keys has height

$$h \leq 2 \log_2(n+1)$$

□ **Proof. INTUITION:**

- Merge their black parents




# Height of a red-black tree

□ **Theorem.** A red-black tree with  $n$  keys has height

$$h \leq 2 \log_2(n+1)$$

□ **Proof. INTUITION:**

- Merge red nodes into their black parents
- This process produces a tree in which each node has 2, 3, or 4 children.
- The 2-3-4 tree has uniform depth  $h'$  of leaves.






# Proof (continued)

- We have  $h' \geq h/2$ , since at most half the leaves on any path are red.
- The number of leaves in each tree is  $n+1$

$$\rightarrow n + 1 \geq 2^{h'}$$

$$\rightarrow \log_2(n + 1) \geq h' \geq h/2$$

$$\rightarrow h \leq 2 \log_2(n + 1)$$



# Query operations

- **Corollary.** The queries SEARCH, MIN, MAX, SUCCESSOR, and PREDECESSOR all run in  $O(\log_2 n)$  time on a red-black tree with  $n$  nodes.




# Modifying operations

- The operations INSERT and DELETE cause modifications to the red-black tree:
  - the operation itself,
  - color changes,
  - restructuring the links of the tree via ***“rotations”***.



# Insertion into a red-black tree

- IDEA: Insert  $x$  in tree. Color  $x$  red. Red-black property 2&4 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.
- Example:
  - Insert  $x = 15$ .




# Insertion into a red-black tree

- IDEA: Insert  $x$  in tree. Color  $x$  red. Red-black property 2&4 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

- Example:

- Insert  $x = 15$ .
- Recolor, moving the violation up the tree.

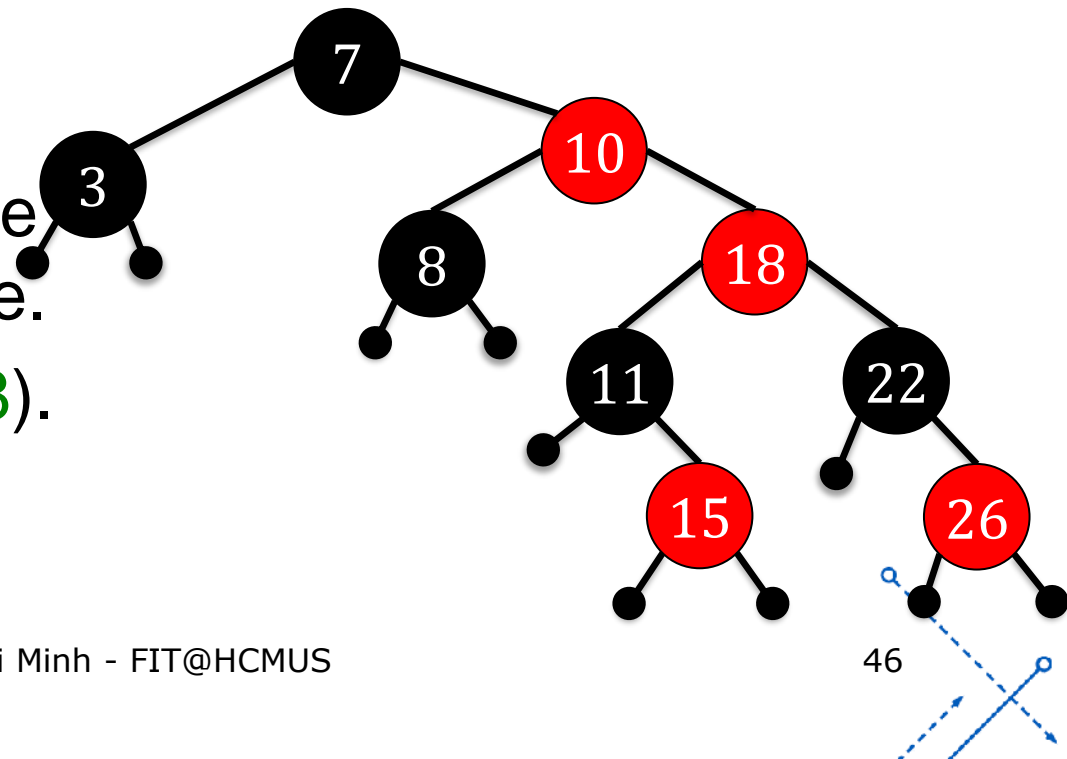


# Insertion into a red-black tree

- IDEA: Insert  $x$  in tree. Color  $x$  red. Red-black property 2&4 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

- Example:

- Insert  $x = 15$ .
- Recolor, moving the violation up the tree.
- RIGHT-ROTATE( $18$ ).




# Insertion into a red-black tree

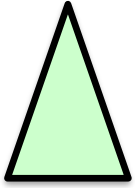
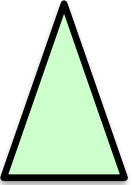
- IDEA: Insert  $x$  in tree. Color  $x$  red. Red-black property 2&4 might be violated. Move the violation up the tree by recoloring until it can be fixed with rotations and recoloring.

## Example:

- Insert  $x = 15$ .
- Recolor, moving the violation up the tree.
- RIGHT-ROTATE(18)
- LEFT-ROTATE(7) and recolor.




# Graphical notation

- Let  denote a subtree with a black root.
- All  's have the same black-height.




# Case 1



(Or, children of  $A$   
are swapped.)


Push  $C$ 's black onto  $A$  and  
 $D$ , and recurse, since  $C$ 's  
parent may be red.

# Case 2



Transform to Case 3.

# Case 3



Done. No more violations of RB properties are possible.

# Analysis

## □ RB-INSERT:

$$O(\log_2 n)$$

- Go up the tree performing Case 1, which only recolors nodes.
- If Case 2 or Case 3 occurs, perform 1 or 2 rotations, and terminate.

## □ RB-DELETE — same asymptotic running time and number of rotations as RB-INSERT (see textbook)

$$O(\log_2 n)$$



# AVL trees vs Red-black trees

	AVL trees	Red-black trees
Balance	More strict	Less strict
Max height	$1.44 \log_2(n+2) - 0.328$	$2 \log_2(n+1)$
INSERT	Slower	Faster
DELETE	Slower	Faster
SEARCH	Faster	Slower
TASKS	Look-up intensive tasks	Insert/delete intensive tasks

# Some applications

---

## □ AVL trees:

- Not much real-life applications.
- Case-study: Documents indexing

## □ Red-black trees:

- Java: `java.util.TreeMap` , `java.util.TreeSet` .
- C++ STL: `map`, `multimap`, `multiset`.
- Linux kernel: completely fair scheduler

# Q&A

