

DATA STRUCTURES & ALGORITHMS

Lecture 6: TREES – Part 3

B-trees, 2-3 trees, 2-3-4 trees

Lecturer: Dr. Nguyen Hai Minh

OUTLINE

□ Balanced Search Tree

- B-tree

- 2-3 tree

- 2-3-4 tree

□ Comparing trees

B-TREES


- Definition
- Example
- Traversal, Search, Insert, Delete
- Advantage of B-trees

Motivation of B-trees

- So far, we have assumed that we can store an entire data structure in *main memory*
- What if we have so much data that it won't fit?
 - Storing it on disk requires different approach to *efficiency*
 - Assuming that a disk spins at 7200 RPM, one revolution occurs in **8.33ms**
 - Crudely speaking, one disk access takes about the same time as **100,000 instructions**


Motivation of B-trees

- Assume that we use an AVL tree to store about 20 million records
 - We end up with a very deep binary tree with lots of different disk accesses; $\log_2 20,000,000 \approx 24$, so this takes about 0.2 seconds
- We know we can't improve on the $\log_2 n$ lower bound on search for a binary tree
- But, the solution is to use **more branches** and thus **reduce the height** of the tree!
 - As **branching increases**, **depth decreases**
 - B-tree: 1970 by Rudolf Bayer & Edward M. McCreight



Definition of a B-tree

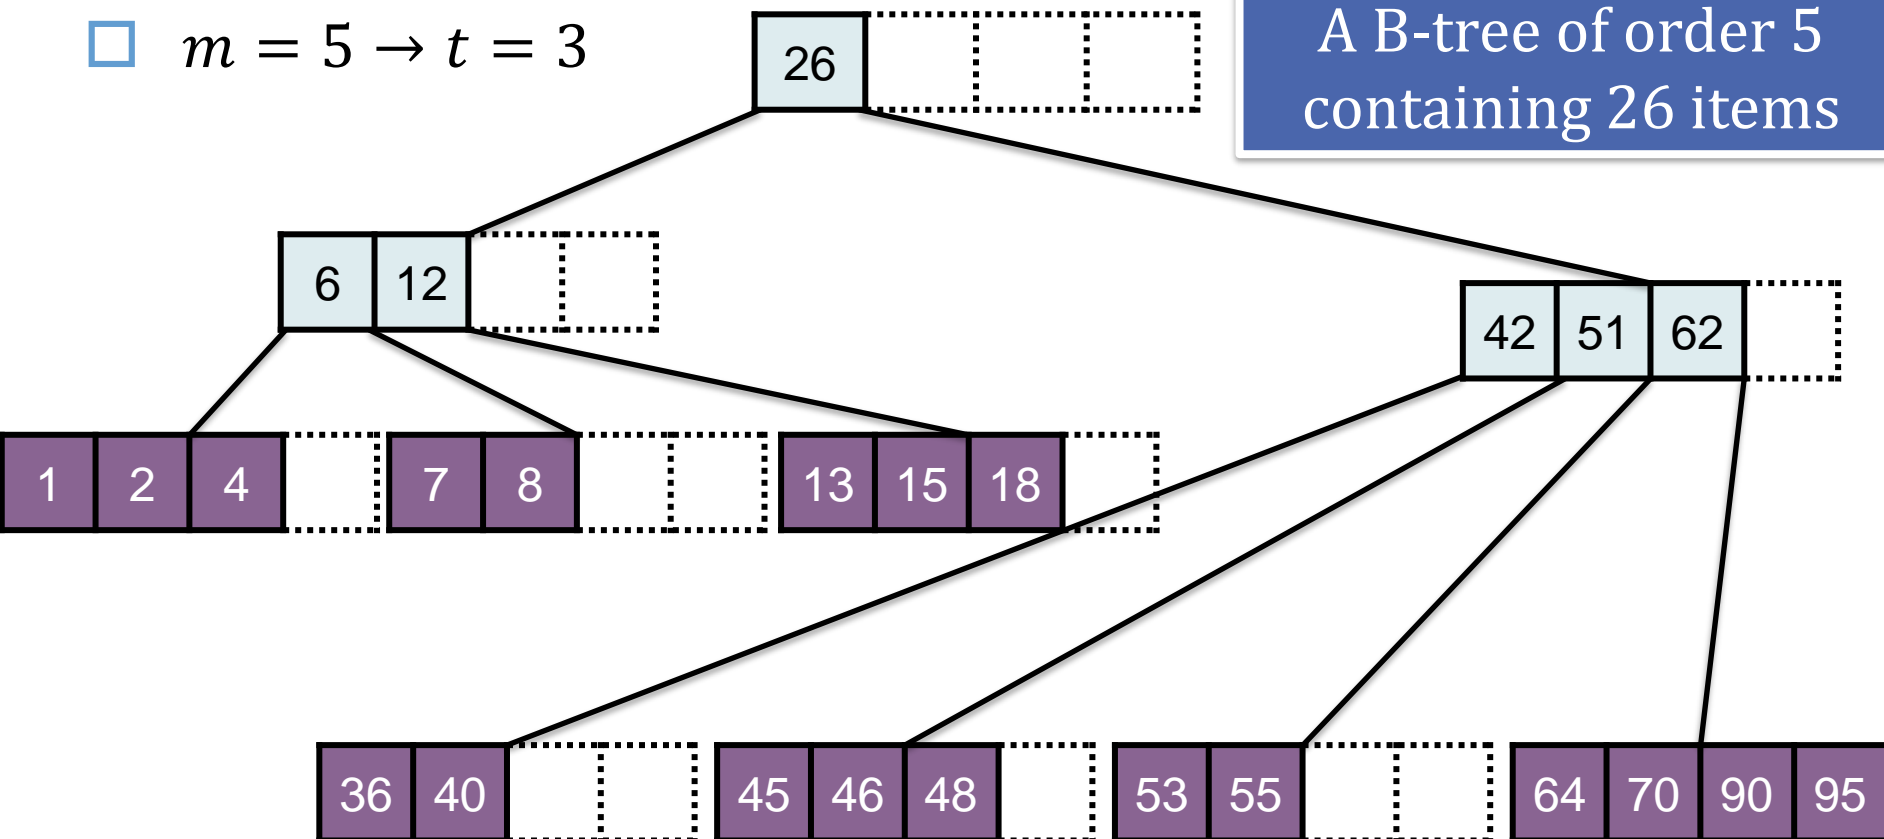
- A B-tree of order m is an m -way tree (i.e., a tree where each node may have up to m children) in which:
 1. The number of keys in each non-leaf node is one less than the number of its children and these keys partition the keys in the children in the fashion of a search tree
 2. All leaves are on the **same level**
 3. All non-leaf nodes except the root have at least $t = \left\lceil \frac{m}{2} \right\rceil$ children
 4. The root is either a leaf node, or it has from 2 to m children
 5. Every nodes except the root has $\left\lceil \frac{m}{2} \right\rceil - 1$ to $m - 1$ keys
- The number m is usually **odd**



An Example B-tree

□ $m = 5 \rightarrow t = 3$

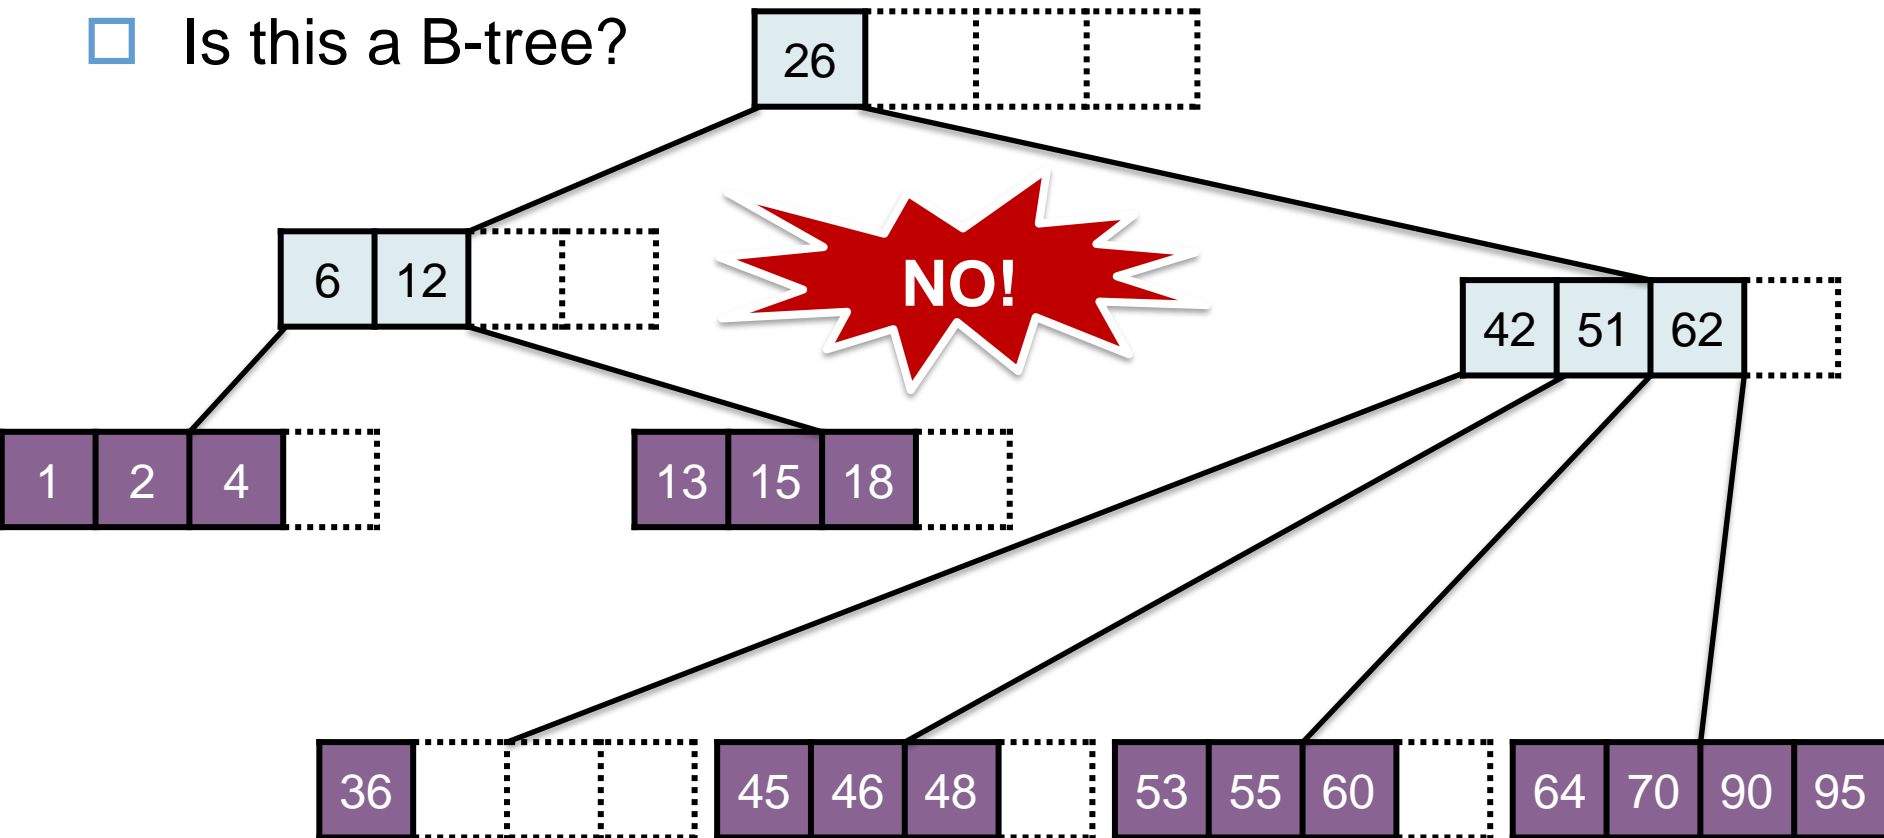
A B-tree of order 5 containing 26 items



- Each internal nodes (except root) has at least 3 children
- Each node has 2-4 keys

An Example B-tree


□ Is this a B-tree?



□ Why?



B-tree Example

- Draw a B-tree of order 3 and height = 3 containing as many keys/data items as possible.
 - How many nodes are there?
 - How many keys does your tree have?



Height of B-tree

- Let h be the height, n is the total keys on a B-tree
- t : minimum degree of m -way B-tree
- Minimum keys on a B-tree:



B-tree's height

- Minimum keys on a B-tree:

$$n \geq 1 + 2(t - 1) + 2t(t - 1) + \dots + 2t^{h-1}(t - 1)$$

$$n \geq 1 + 2(t - 1) \sum_{i=1}^h t^{i-1} = 1 + 2(t - 1) \left(\frac{t^h - 1}{t - 1} \right)$$

$$n \geq 2t^h - 1$$

- **Theorem:** If $n \geq 1$, then for any n -keys B-tree of height h and minimum degree $t \geq 2$,

$$h \leq \log_t \frac{n + 1}{2}$$

$$\rightarrow h = O(\log_t n)$$



B-tree's maximum nodes

- The **maximum** number of keys in a B-tree of order m and height h :

root $m - 1$

level 1 $m(m - 1)$

level 2 $m^2(m - 1)$

...

level h $m^h(m - 1)$

- So, the total number of keys is:

$$n \leq (1 + m + m^2 + m^3 + \dots + m^h)(m - 1) =$$

$$\left[\frac{m^{h+1} - 1}{m - 1} \right] (m - 1) = m^{h+1} - 1$$

$$\rightarrow h \geq \log_m(n + 1) - 1$$



Traversal and Search in B-tree

□ Traversal in B-tree:

- Similar to *In-Order Traversal* of Binary Tree.

□ Search in B-tree:

- Similar to searching in a Binary Search Tree, except that instead of making a binary, or “two-way,” branching decision at each node, we make a *multiway* branching decision according to the number of the node’s children.



B-tree-SEARCH

B-tree-SEARCH(x, k)

$$O(th) = O(t \log_t n)$$

// x : pointer to the root node of a subtree, k : key to be searched

1 $i = 0$

2 **while** $i < x.n$ and $k > x.key_i$

3 $i = i + 1$

$$O(t)$$

4 **if** $i < x.n$ and $k == x.key_i$

5 **return** (x, i)

6 **elseif** $x.leaf$

7 **return** NULL

8 **else** DISK-READ($x.c_i$)

9 **return** B-tree-SEARCH($x.c_i, k$)

$$O(h)$$

Inserting into a B-tree

$$O(t \log_t n)$$

- Attempt to insert the new key into a leaf
- If *leaf* becomes *too big*,
 - Split the leaf into two
 - Promoting the middle key to the leaf's parent
- If *the parent* becomes *too big*
 - Split the parent into two
 - Promoting the middle key
- This strategy might have to be repeated all the way to the top.
- If necessary, the root is split in two and the middle key is promoted to a new root, *making the tree one level higher*



Inserting into a B-tree – Example

- Suppose we start with an empty B-tree and keys arrive in the following order:

1 12 8 2 25 5 14 28 17 7 52 16 48 68 3
26 29 53 55 45


- We want to construct a B-tree of order 5



Inserting into a B-tree – Example



➤ 1 12 8 2

1	2	8	12
---	---	---	----




Inserting into a B-tree – Example

➤ 1 12 8 2 **25**




Inserting into a B-tree – Example

➤ 1 12 8 2 25 **5 14 28**





Inserting into a B-tree – Example

➤ 1 12 8 2 25 5 14 28 **17**





Inserting into a B-tree – Example

➤ 1 12 8 2 25 5 14 28 17 **7 52 16 48**





Inserting into a B-tree – Example

➤ 1 12 8 2 25 5 14 28 17 7 52 16 48 **68**




Inserting into a B-tree – Example

➤ 1 12 8 2 25 5 14 28 17 7 52 16 48 68 **3**





Inserting into a B-tree – Example

➤ 1 12 8 2 25 5 14 28 17 7 52 16 48 68 3 **26 29 53 55**



Inserting into a B-tree – Example

➤ 1 12 8 2 25 5 14 28 17 7 52 16 48 68 3 26 29 53 55 **45**





Insert 45



Inserting into a B-tree – Example


➤ 1 12 8 2 25 5 14 28 17 7 52 16 48 68 3 26 29 53 55 **45**



Removal from a B-tree

□ Remove a key k

1. If k is in a **leaf** node, and removing it doesn't cause that leaf node to have too few keys, then simply remove k .
2. If k is **NOT** in a leaf then it is guaranteed (by the nature of a B-tree) that its predecessor or successor will be in a leaf -- in this case we can delete k and promote the **predecessor** or **successor** of k to k 's position.



Removal from a B-tree (2)


- (1) & (2) may lead to a leaf node L has less than min. number of keys
- Look at the siblings immediately adjacent to the leaf
 - 3: if one of them has more than the min. number of keys then we can promote one of its keys to the parent and take the parent key into L
 - 4: otherwise, combine L and one of its neighbours with their shared parent, repeat the process up to the root, if required

$$O(t \log_t n)$$



Type #1: Simple leaf deletion

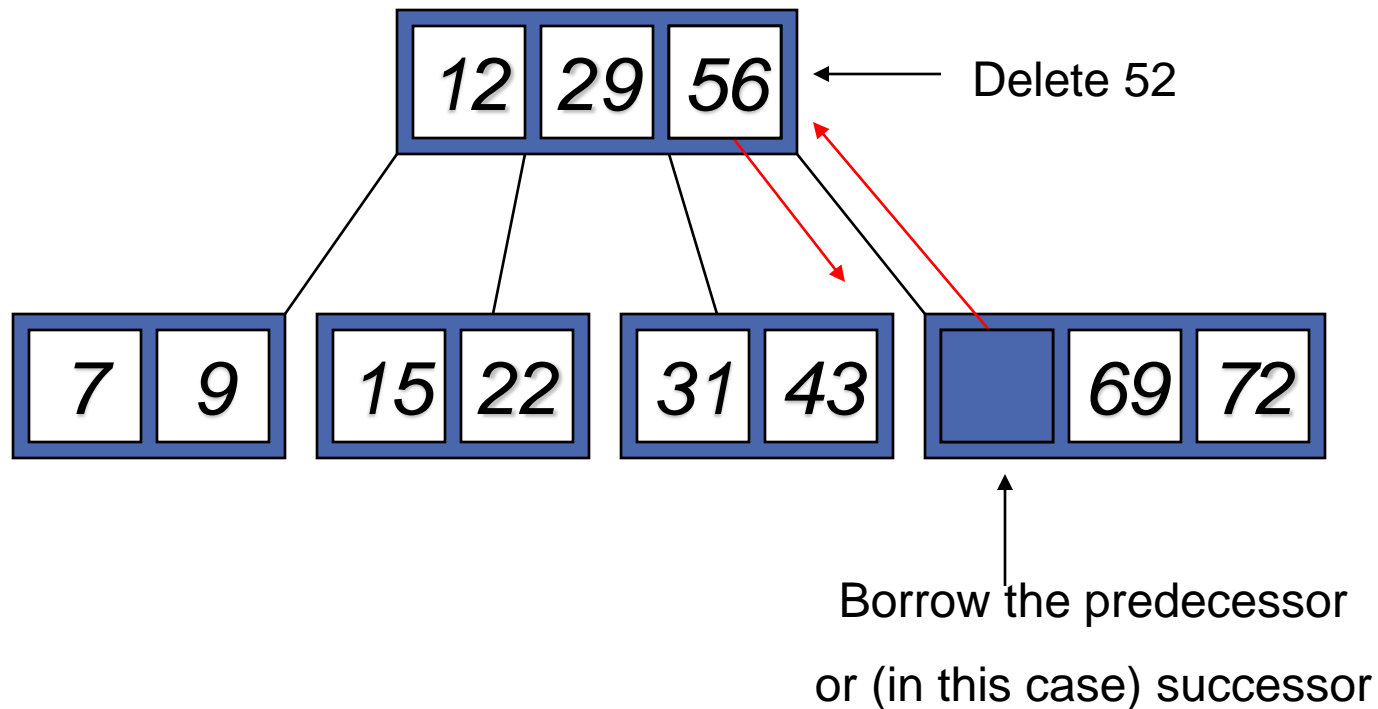
Assuming a 5-way
B-tree, as before...



Delete 2: Since there are enough
keys in the node, just delete it


Note when printed: this slide is animated

Type #2: Simple non-leaf deletion




Note when printed: this slide is animated

Type #4: Too few keys in node and its siblings




Note when printed: this slide is animated

Type #4: Too few keys in node and its siblings




Note when printed: this slide is animated

Type #3: Enough siblings



Note when printed: this slide is animated

Type #3: Enough siblings



Note when printed: this slide is animated

Exercise in B-tree

- Given 5-way B-tree created by these data:
3, 7, 9, 23, 45, 1, 5, 14, 25, 24, 13, 11, 8, 19, 4,
31, 35, 56
- Add these further keys: 2, 6, 12
- Delete these keys: 4, 5, 7, 3, 14



Reasons for using B-trees

- When searching tables held on disc, the cost of each disc transfer is high but doesn't depend much on the amount of data transferred, especially if consecutive items are transferred
 - If we use a B-tree of order 101, say, we can transfer each node in one disc read operation
 - A B-tree of order 101 and height 3 can hold $101^4 - 1$ items (approximately 100 million) and any item can be accessed with 3 disc reads (assuming we hold the root in memory)
 - B-trees are always balanced (since the leaves are all at the same level) → smallest height is guarantee



2-3 TREES

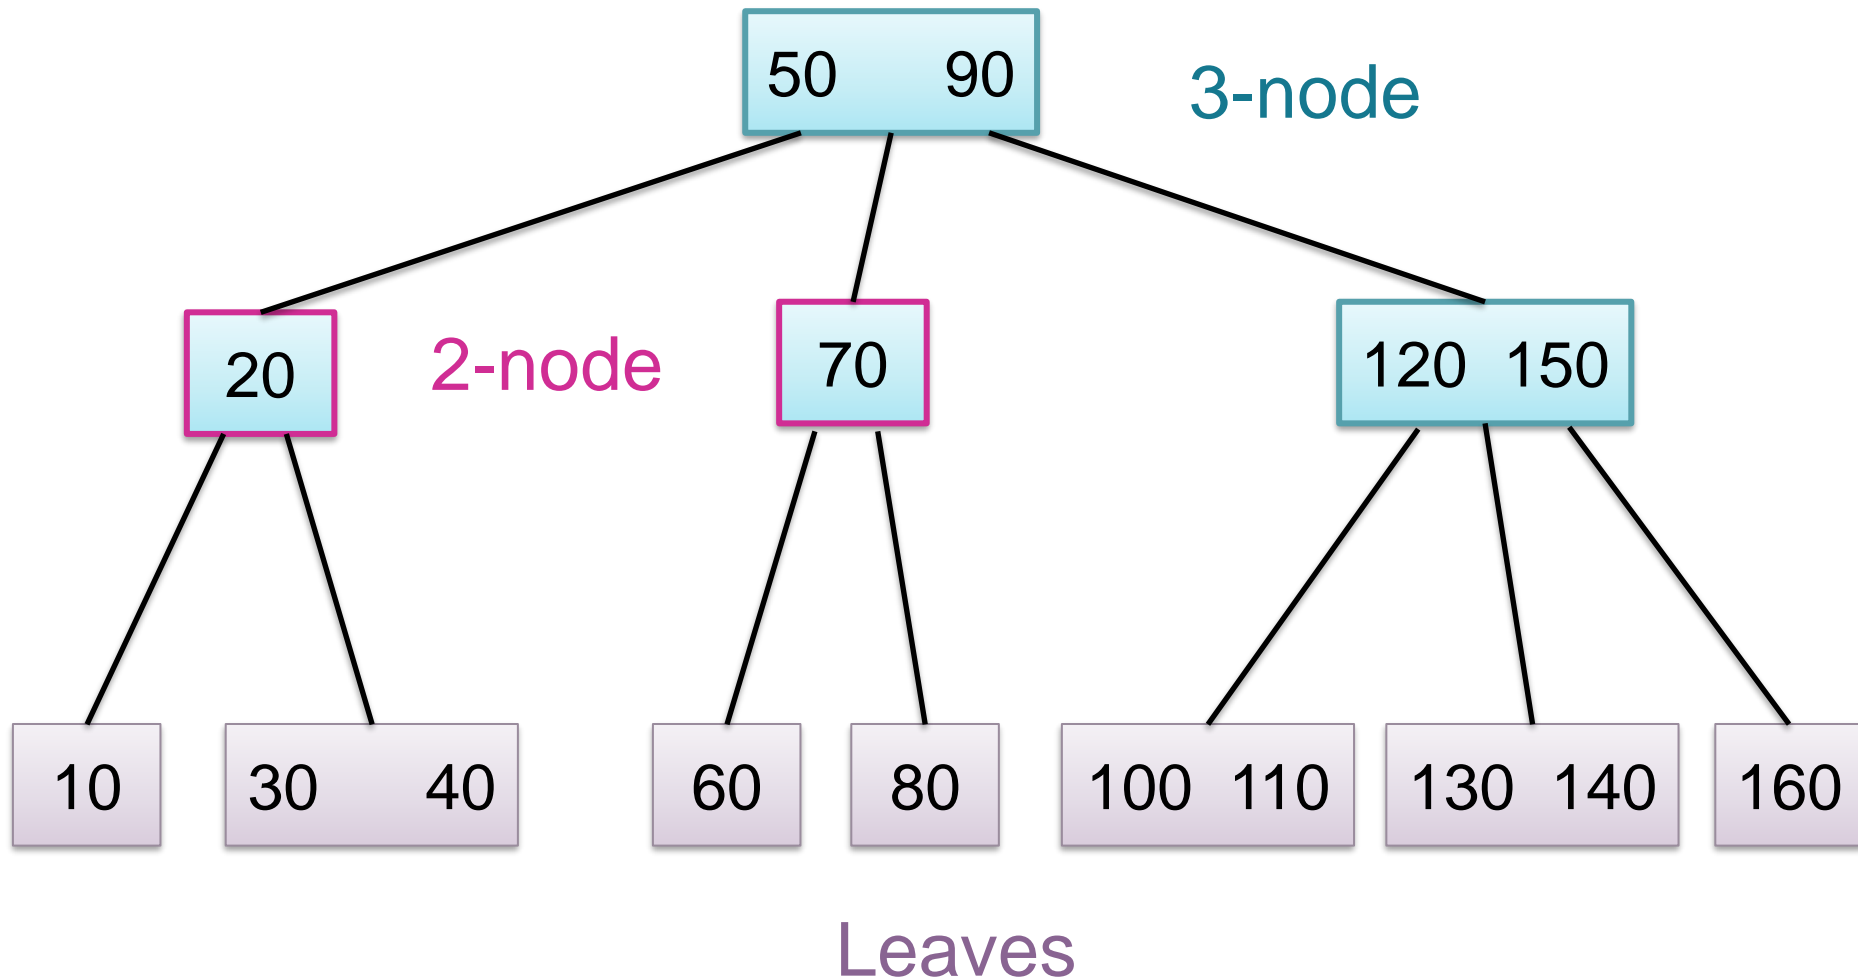
- Definition
- Traversal & Searching
- Insert & Delete
- Why 2-3 tree?

Definition of 2-3 tree

- For B-tree, if we take $m = 3$, we get a **2-3 tree**, in which non-leaf nodes have two or three children (i.e., one or two keys)
- **Definition:** A **2-3 tree** is a B-tree in which:
 - Each internal node has either **2** or **3** children.
 - A node with **2 children** is called a **2-node**.
 - A node with **3 children** is called a **3-node**.
 - All leaves are at the **same level** (leaf contains 1 or 2 keys)



2-3 tree – Example



Height of a 2-3 tree

- A 2-3 tree is **not** a binary tree.
 - If a 2-3 tree contains only 2 nodes, it is likely a perfect binary tree $\rightarrow n = 2^{h+1} - 1$
 - If, on the other hand, some of the internal nodes of a 2-3 tree do have 3 children, the tree will contain more nodes than a perfect binary tree of the same height $\rightarrow n > 2^{h+1} - 1$
- Therefore, height of a 2-3 tree is:


$$h \leq \lceil \log_2(n + 1) - 1 \rceil$$

$$O(\log_2 n)$$



Traversing a 2-3 tree

- You can traverse a 2-3 tree in sorted order by performing the analogue of an in-order traversal on a binary tree





Searching a 2-3 tree

- The ordering of keys in a 2-3 tree is analogous to the ordering for a BST
 - Searching on a 2-3 tree is efficient & quite similar to a BST
- The same efficiency:
 - A BST with n nodes cannot be shorter than $\lceil \log_2(n + 1) - 1 \rceil$
 - A 2-3 tree with n nodes cannot be taller than $\lceil \log_2(n + 1) - 1 \rceil$
 - A node in a 2-3 tree has at most 2 keys.
- *Then, why should we use 2-3 tree?*



Inserting to a 2-3 tree


- A balanced BST and a 2-3 tree with the same keys



- Now, insert 39, 38, ..., 32 to these trees


Inserting to a 2-3 tree

- Now, insert 39, 38, ..., 32 to these trees
- Binary Search Tree
 - Quickly loose its balance




Inserting to a 2-3 tree

- Now, insert 39, 38, ..., 32 to these trees
- 2-3 tree:
 - Retains its structure





Inserting to a 2-3 tree

□ Insert 39




Inserting to a 2-3 tree

□ Insert 38




Inserting to a 2-3 tree

□ Insert 38






Inserting to a 2-3 tree

□ Insert 37




Inserting to a 2-3 tree

□ Insert 36



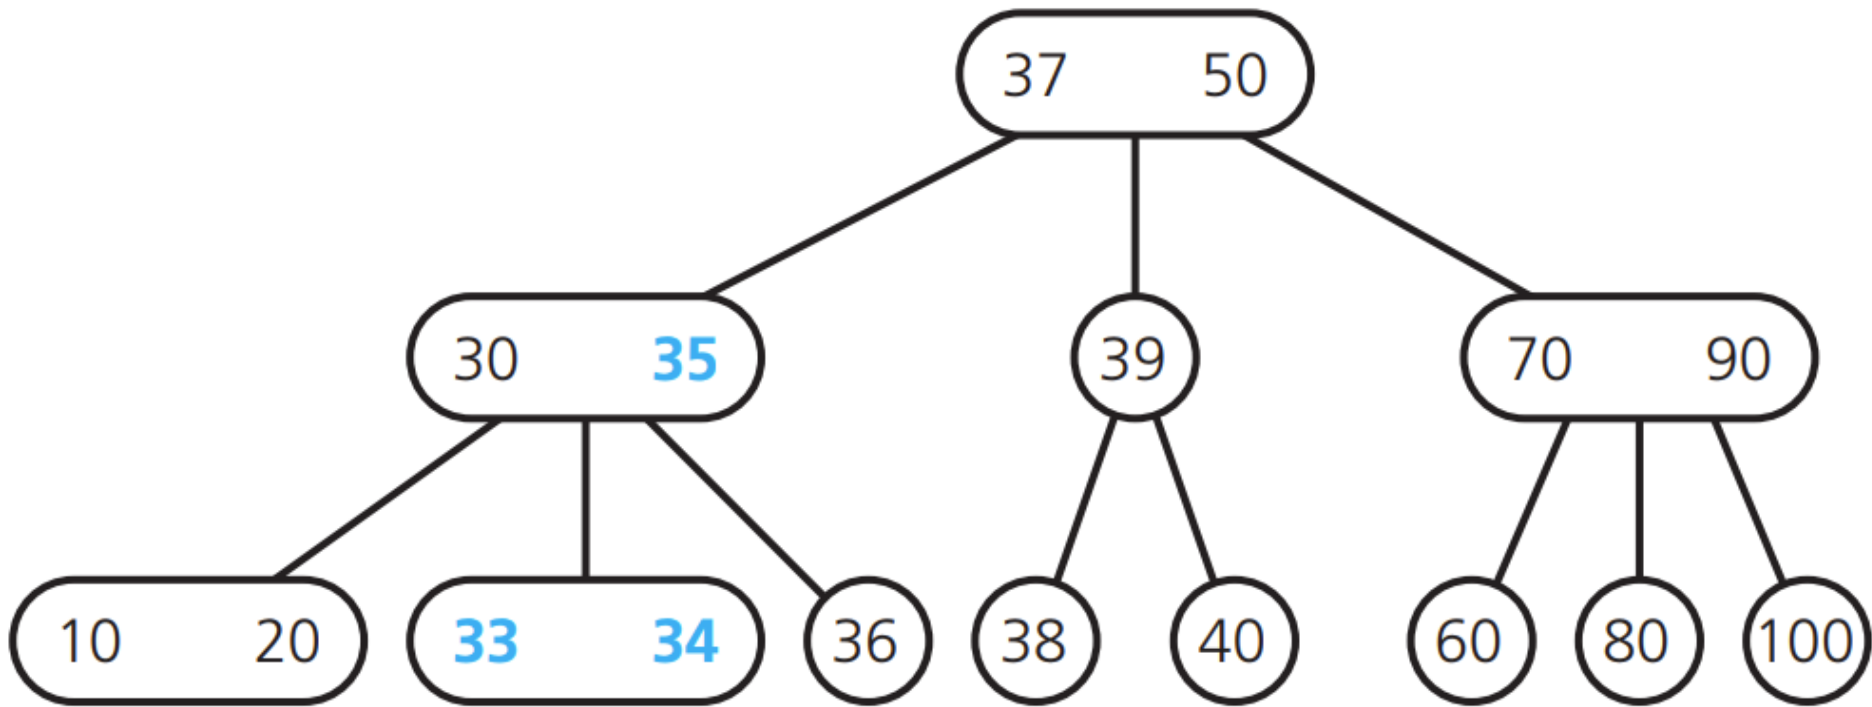
Inserting to a 2-3 tree

□ Insert 36




Inserting to a 2-3 tree

□ Insert 35, 34, 33



Deleting from a 2-3 tree


□ Delete 70:



Swap with inorder successor


Deleting from a 2-3 tree

□ Delete 70:




Deleting from a 2-3 tree


□ Delete 70:




Delete value from leaf



Merge nodes by deleting empty leaf and moving 80 down



Result tree




Deleting from a 2-3 tree

□ Delete 70, 100, 80:


B-tree

(a)



Binary Search Tree

(b)



Again, Why 2-3 tree?

- Searching: sometimes you have to make 2 comparisons to get pass a 3-node.
 - However, it is still $O(\log n)$ running time
- Insertion/Deletion needs extra work: split nodes, merge nodes.
 - However, this extra work is not a real concern
 - It is easier to keep the tree balanced than a normal BST.



2-3-4 TREES

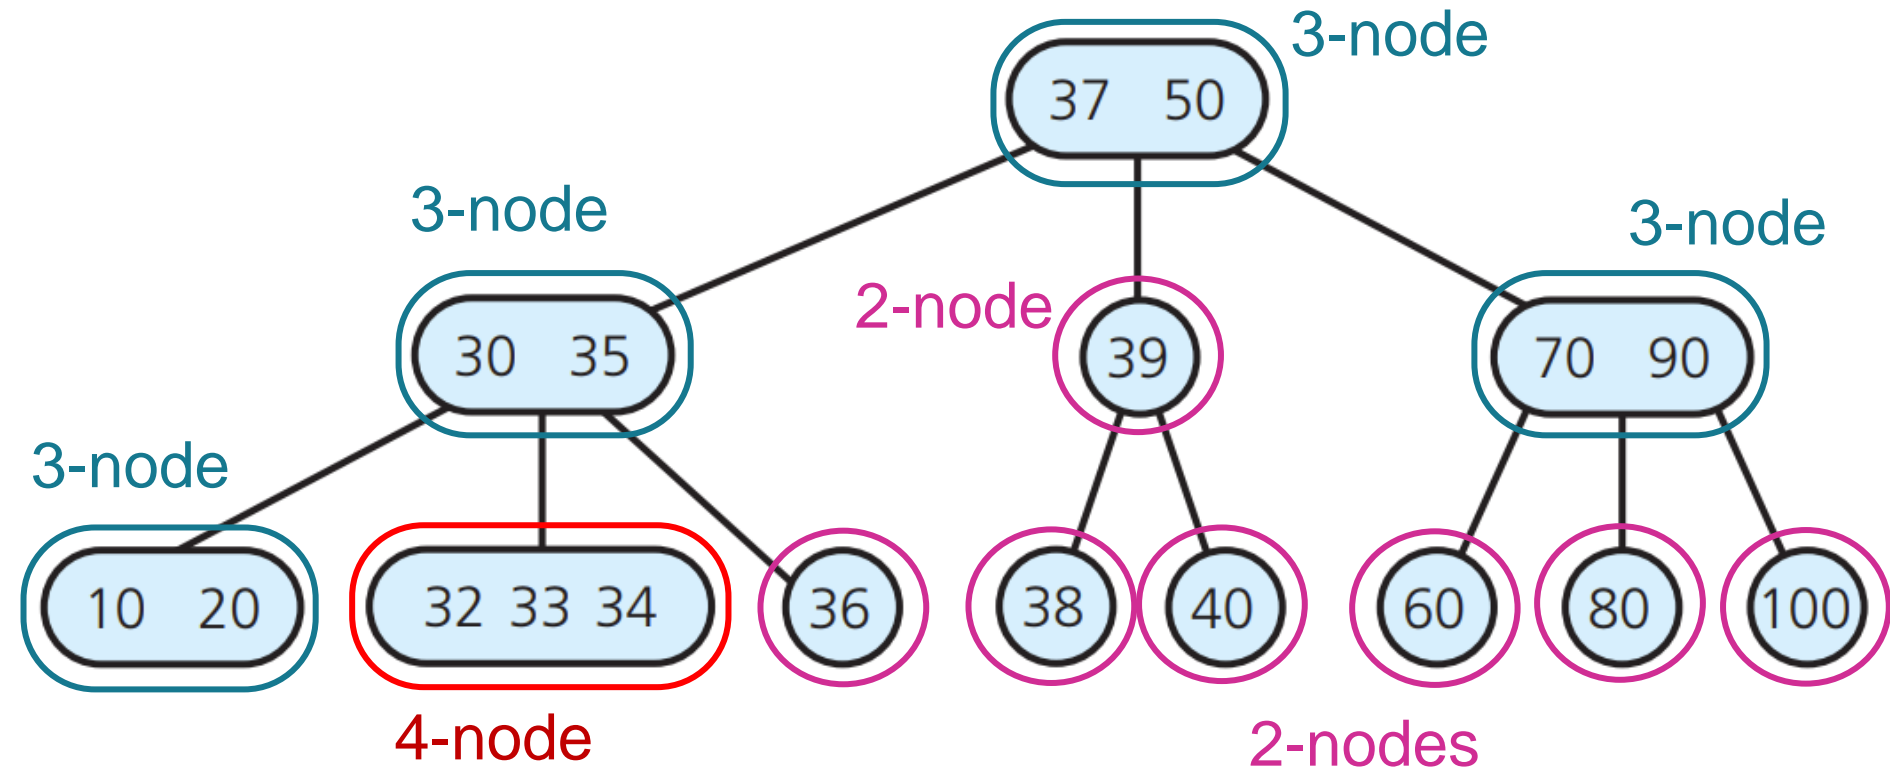
Definition of 2-3-4 tree

- If a 2-3 tree is good, are trees with more than 3 children better?
- **Definition:** A 2-3-4 tree is a B-tree in which:
 - Each internal node has either 2, 3 or 4 children.
 - 2-node: node contains 1 key
 - 3-node: node contains 2 keys
 - 4-node: node contains 3 keys
 - All leaves are at the same level, each contains 1 to 3 keys

*Different with
2-3 tree*

Example of 2-3-4 tree

- 2-3-4 tree of the same data as the previous 2-3 tree



Searching and Traversing a 2-3-4 tree

- Traversing & Searching are the same as in B-tree (or 2-3 tree)
- Example:
 - Search for 31 on the Example 2-3-4 tree:
 - Search the left subtree since $31 < 37$
 - Search the middle subtree of the node $\langle 30 \ 35 \rangle$ since $31 > 30$ and $31 < 35$
 - Terminate the search at the left child pointer of $\langle 32 \ 33 \ 34 \rangle$ since $31 < 32$
 - Result: 31 is not in the tree.




Inserting to 2-3-4 tree

- For inserting a key to a 2-3 tree, the search algorithm traces a path from the root to a leaf and then backs up from the leaf as it splits nodes.
- For a 2-3-4 tree: **splits 4-node** as soon as it encounters them on the way down from the root to the leaf. → *Avoid return path*
 - Each 4-node either:
 1. Be the root, or
 2. Have a 2-node parent, or
 3. Have a 3-node parent



Splitting a 4-node in 2-3-4 tree


1. The 4-node is the root



Splitting a 4-node in 2-3-4 tree

2. The 4-node has a 2-node parent:

- 4-node is left child:




- 4 node is right child?

Splitting a 4-node in 2-3-4 tree

3. The 4-node has a 3-node parent


- 4-node is left child:




- 4-node is middle child?
- 4-node is right child?

Inserting to 2-3-4 tree


- Insert **20** into a one-node 2-3-4 tree



- This is a 4-node → split it




- Then, insert to the leaf node




Inserting to 2-3-4 tree

- Insert **50, 40** do not require split nodes




- Insert **70** → split the $\langle 40\ 50\ 60 \rangle$ node, then insert




Inserting to 2-3-4 tree

□ Insert **80, 15** do not require split nodes





□ Insert **90** → split the <60 70 80> node, then insert



Inserting to 2-3-4 tree

□ Insert **100** → split the $\langle 30 \ 50 \ 70 \rangle$ node, then insert



Deleting from a 2-3-4 tree

- Similar to deleting a key from a 2-3 tree:
 - Locate the node i that contains the key k to remove
 - Find k 's successor (at leaves) and swap with k
 - Ensure that k doesn't occur in a 2-node so that we can perform one pass removal (unlike removal from a 2-3 tree)
- *Transform each 2-node that we encounter during the search for k into a 3-node or 4-node*
- Detail of deleting from a 2-3-4 are left as exercises



Conclusion

□ Pros:

- A 2-3-4 is balanced
- Its insertion and removal operations use only one pass from root to leaf.

□ Cons:

- Requires more storage than a binary search tree that contains the same data
- Red-black tree can represent a 2-3-4 tree that retains the advantages of a 2-3-4 tree without the storage overhead.



Comparing Trees

- ❑ **Binary search tree** can become *unbalanced* and *lose* their good time complexity
- ❑ A **2-3 tree** and a **2-3-4 tree** are variants of a BST that keeps the tree balanced easily.
- ❑ A **Red-black tree** is a binary tree representation of a 2-3-4 tree that requires less storage than a 2-3-4 tree. Insertions and removals in RB tree are more efficient than 2-3-4 tree.
- ❑ **AVL tree** is strict binary trees that *overcome the balance problem*
- ❑ **B-tree** is balanced search tree designed to work well on disk drivers or other direct-access secondary storage devices



What's next?

□ After today:

■ Reading

- ✓ B-tree: Textbook 1 chap 18 (page 655~)
- ✓ 2-3 tree: Textbook 2 sec 19.2 (page 569~)
- ✓ 2-3-4 tree: Textbook 2 sec 19.3 (page 585~)
- ✓ 2-3-4 tree to Red-black tree: Textbook 2 sec 19.4 (page 592~)

■ Do homework 6.2

□ Next class (August 2nd):

- Individual Assignment 4
- Lecture 7: Graphs



Q&A